

- ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022. URL <https://arxiv.org/abs/2212.08073>.
- [5] Nick Bostrom. *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, 2014.
 - [6] Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nicholas L. Turner, Cem Anil, Carson Denison, Amanda Askell, et al. Towards monosemanticity: Decomposing language models with dictionary learning. 2023. URL <https://transformer-circuits.pub/2023/monosemantic-features>. Transformer Circuits (online publication).
 - [7] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt. Discovering latent knowledge in language models without supervision. In *International Conference on Learning Representations (ICLR)*, 2023. Originally arXiv 2022. Proposes finding "truth" vectors inside the model.
 - [8] Stephen Casper et al. Explore, establish, exploit: Red teaming as a search problem. In *Proceedings of the 2023 Conference on Neural Information Processing Systems (NeurIPS)*, 2023.
 - [9] Stephen Casper et al. Explore, establish, exploit: Red teaming as a search problem. In *Proceedings of the 2023 Conference on Neural Information Processing Systems (NeurIPS)*, 2023. Discusses the non-exhaustive nature of adversarial testing.
 - [10] Gelei Deng, Yi Liu, Yuekang Li, Kailong Wang, Ying Zhang, Zefeng Li, Haoyu Wang, Tianwei Zhang, and Yang Liu. Masterkey: Automated jailbreaking of large language model chatbots. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*. Internet Society, 2024. doi: 10.14722/ndss.2024.24188. URL <https://www.ndss-symposium.org/wp-content/uploads/2024-188-paper.pdf>.
 - [11] Yue Deng, Wenxuan Zhang, Sinno Jialin Pan, and Lidong Bing. Multilingual jailbreak challenges in large language models. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://www.cse.cuhk.edu.hk/~sinnopan/publications/%5BICLR24%5DMultilingual%20Jailbreak%20Challenges%20in%20Large%20Language%20Models.pdf>.
 - [12] Pierre Duhem. *The aim and structure of physical theory*. Princeton University Press, 1954. Classic formulation of the impossibility of crucial experiments.
 - [13] Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, et al. Toy models of superposition. *arXiv preprint arXiv:2209.10652*, 2022. URL <https://arxiv.org/abs/2209.10652>.
 - [14] Tom Everitt, Victoria Krakovna, Laurent Orseau, and Shane Legg. Reinforcement learning with a corrupted reward channel. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI)*, 2017.
 - [15] Tom Everitt, Victoria Krakovna, and Laurent Orseau. Reward tampering problems and solutions in reinforcement learning. *Journal of Artificial Intelligence Research*, 72:101–165, 2021.
 - [16] Tom Everitt et al. Reinforcement learning with a corrupted reward channel. *IJCAI*, 2017.

- [17] Sebastian Gehrmann et al. The gem benchmark: Natural language generation, its evaluation and metrics. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2022.
- [18] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A Wichmann. Shortcut learning in deep neural networks. In *Nature Machine Intelligence*, 2020. Empirical analogue to our theoretical evaluation awareness.
- [19] Helia Hashemi, Jason Eisner, et al. Llm-rubric: A multidimensional, calibrated approach to automatic evaluation of natural language texts. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [20] Dan Hendrycks, Kevin Zhao, Steven Basart, Jacob Steinhardt, and Dawn Song. Natural adversarial examples. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021.
- [21] Evan Hubinger et al. Risks from learned optimization in advanced machine learning systems. *arXiv preprint arXiv:1906.01820*, 2019.
- [22] Geoffrey Irving, Paul Christiano, and Dario Amodei. Ai safety via debate. *arXiv preprint arXiv:1805.00899*, 2018. URL <https://arxiv.org/abs/1805.00899>.
- [23] Jiaming Ji, Manya Liu, Josef Dai, Xuehai Pan, Chi Zhang, Ce Bian, Boyuan Chen, Ruiyang Sun, Yizhou Wang, and Yaodong Yang. Beavertails: Towards improved safety alignment of llm via a human-preference dataset. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- [24] Raymond Koo, Mina Lee, et al. Benchmarking cognitive biases in large language models as evaluators. In *Findings of the Association for Computational Linguistics (ACL)*, 2024.
- [25] Victoria Krakovna et al. Specification gaming: the flip side of ai ingenuity. *DeepMind Blog*, 2020.
- [26] Michael Krumdick, Charles Lovering, and Varshini Reddy. No free labels: Limitations of llm-as-a-judge without human grounding. In *Findings of the Association for Computational Linguistics (ACL)*, 2024.
- [27] Md Tahmid Rahman Laskar, Sawsan Alqahtani, et al. A systematic survey and critical review on evaluating large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.
- [28] Jan Leike, David Krueger, Tom Everitt, Miljan Martic, Andrew Lefrancq, Laurent Orseau, Shane Legg, and Dario Amodei. Scalable agent alignment via reward modeling: a research direction. *arXiv preprint arXiv:1811.07871*, 2018. URL <https://arxiv.org/abs/1811.07871>.
- [29] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*, 2022. URL <https://arxiv.org/abs/2211.09110>.
- [30] David Manheim and Scott Garrabrant. Categorizing variants of goodhart’s law. *arXiv preprint arXiv:1803.04585*, 2018. Foundational theory on proxy optimization.

- [31] Anay Mehrotra, Manolis Zampetakis, Paul Kassianik, Blaine Nelson, Hyrum Anderson, Yaron Singer, and Amin Karbasi. Tree of attacks: Jailbreaking black-box LLMs automatically. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/70702e8cbb4890b4a467b984ae59828a-Paper-Conference.pdf.
- [32] Andrew Y. Ng and Stuart J. Russell. Algorithms for inverse reinforcement learning. In *Proceedings of the Seventeenth International Conference on Machine Learning (ICML)*, 2000.
- [33] Richard Ngo, Lawrence Chan, and Sören Mindermann. The alignment problem from a deep learning perspective. *arXiv preprint arXiv:2209.00626*, 2022.
- [34] Ethan Perez et al. Discovering language model behaviors with model-written evaluations. *arXiv preprint arXiv:2212.09251*, 2022.
- [35] Willard V Quine. Two dogmas of empiricism. *The philosophical review*, 60(1):20–43, 1951. Foundational text on the underdetermination of theory by evidence.
- [36] Ibne Farabi Shihab, Sanjeda Akter, and Anuj Sharma. Detecting proxy gaming in rl and llm alignment via evaluator stress tests. *arXiv preprint arXiv:2507.05619*, 2025.
- [37] Ibne Farabi Shihab, Sanjeda Akter, and Anuj Sharma. Detecting proxy gaming in rl and llm alignment via evaluator stress tests. *arXiv preprint arXiv:2507.05619*, 2025. Included as supporting empirical evidence.
- [38] Aarohi Srivastava et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint arXiv:2206.04615*, 2022. URL <https://arxiv.org/abs/2206.04615>.
- [39] Teun van der Weij, Felix Hofstätter, Oliver Jaffe, Samuel F. Brown, and Francis Rhys Ward. Ai sandbagging: Language models can strategically underperform on evaluations. *arXiv preprint arXiv:2406.07358*, 2024. Included as recent empirical evidence.
- [40] Yuntao Wang, Yushi Bai, Zhen Wang, Zhilin Yang, Yang Zhu, Wen Wang, Jian Li, Junxian Wang, Wei Li, and Zhi-Hua Zhou. Pku-saferlhf: A safety benchmark for reinforcement learning from human feedback. *arXiv preprint arXiv:2307.02483*, 2023.
- [41] Thomas Wolf et al. Fundamental limits of alignment in large language models. *Journal of Artificial Intelligence Research*, 79:412–445, 2024. Explores the information-theoretic boundaries of alignment.
- [42] Jinchuan Zhang, Yan Zhou, Yaxin Liu, Ziming Li, and Songlin Hu. Holistic automated red teaming for large language models through top-down test case generation and multi-turn interaction. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 13711–13736. Association for Computational Linguistics, 2024. URL <https://aclanthology.org/2024.emnlp-main.760/>.
- [43] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, et al. Representation engineering: A top-down approach to ai transparency. *arXiv preprint arXiv:2310.01405*, 2023. Key paper on manipulating internal states for safety.